

CSCI-UA.0480-003
Parallel Computing
Lab Assignment 1

In this lab you will implement a method for solving a group of linear equations using MPI.

What will your program do?

Given a set of n equations with n unknowns (x_1 to x_n), your program will calculate the values of x_1 to x_n within an error margin of $e\%$.

The format of the file is:

- line1: #unknowns
- line2: absolute relative error
- Initial values for each unknown
- line 3 till end: the coefficients for each equation. Each equation on a line. On the same line and after all the coefficients you will find the constant of the corresponding equation.

For example, if we want to solve a system of 3 linear equations, you can have a file like this one:

```
3
0.01
2 3 4
5 1 3 6
3 7 2 8
3 6 9 6
```

The above file corresponds to the following set of equation:

$$\begin{aligned} 5X_1 + X_2 + 3X_3 &= 6 \\ 3X_1 + 7X_2 + 2X_3 &= 8 \\ 3X_1 + 6X_2 + 9X_3 &= 6 \end{aligned}$$

The third line in the file tells us that the initial values for X_1 is 2, for X_2 is 3, and for X_3 is 4. Those values may not be the solution, or are very far from the solution that must be within 1% of the real values (as given by the 0.01 in line 2).

How will your program do that?

We start with a set of n equations and n unknowns, like this:

$$a_{11}x_1 + a_{12}x_2 + a_{13}x_3 + \dots + a_{1n}x_n = b_1$$

$$a_{21}x_1 + a_{22}x_2 + a_{23}x_3 + \dots + a_{2n}x_n = b_2$$

$$\vdots$$

$$a_{n1}x_1 + a_{n2}x_2 + a_{n3}x_3 + \dots + a_{nn}x_n = b_n$$

You are given all a_{ij} and b_1 to b_n . You need to calculate all Xs.

Here are the steps:

1. Rewrite each equation such has the left-hand-side is one of the unknowns.

Rewriting each equation

$$x_1 = \frac{c_1 - a_{12}x_2 - a_{13}x_3 - \dots - a_{1n}x_n}{a_{11}} \quad \leftarrow \text{From Equation 1}$$

$$x_2 = \frac{c_2 - a_{21}x_1 - a_{23}x_3 - \dots - a_{2n}x_n}{a_{22}} \quad \leftarrow \text{From equation 2}$$

$$\vdots$$

$$x_{n-1} = \frac{c_{n-1} - a_{n-1,1}x_1 - a_{n-1,2}x_2 - \dots - a_{n-1,n-2}x_{n-2} - a_{n-1,n}x_n}{a_{n-1,n-1}} \quad \leftarrow \text{From equation n-1}$$

$$x_n = \frac{c_n - a_{n1}x_1 - a_{n2}x_2 - \dots - a_{nn-1}x_{n-1}}{a_{nn}} \quad \leftarrow \text{From equation n}$$

Note: The Cs above refer to the constants, which are the b_1 to b_n .

In general:

$$x_i = \frac{c_i - \sum_{\substack{j=1 \\ j \neq i}}^n a_{ij}x_j}{a_{ii}}, i = 1, 2, \dots, n.$$

2. Remember that you were also given some initial values for the Xs in the input file. The absolute relative error is:

$$\left| \epsilon_a \right|_i = \left| \frac{x_i^{new} - x_i^{old}}{x_i^{new}} \right| \times 100$$

Therefore, our goal is to reduce absolute relative error for each unknown to make it less or equal to relative error given in the input file (2nd line). Note: You need to multiply the error given in the file by 100 to match it with the above equation, or to not multiply the above equation by 100.

3. Substitute the initial values in the equation of each X_i to get new value for X_i . Now we have a new set of Xs.
4. Calculate the absolute relative errors for each X.
5. If all errors are equal or less the given number (2nd line in the file) then you are done.

6. Otherwise go back to step 3 with the set of new Xs as X_{old} .

What is the input to your program?

The input to your program is a text file named xxxx.eq where xxxx can be any name. We already discussed the file format.

What is the output of your program?

Your program must output to stdout (the screen) the value of each unknown. The output must look like:

2
3
4

Where 2 correspond to the value of X_1 , 3 corresponds to X_2 , and 4 corresponds to X_3, \dots .

In the last line of the output show the number of iterations as:

total number of iterations: 5

What do I do after I finish my program?

We have provided you with a reference program *gsref* so you can check the correctness of your code. We will test your submission against this reference, but with other input files than the distributed ones.

After you finish the parallel version of your program, compile it with:

mpicc -o gs gs.c

Where gs.c is your code. We provide a skeleton file to help you start.

After you compile your program and check its correctness do the following:

- Measure the time of your program (using *time* command) for 1, 2, 4, 8, 16, 32, and 64 processes. You may need to measure the time few times (~ 5) for each and take the average. Do the following with *small.txt*, *medium.txt*, *large.txt*, and *huge.txt*. Note that sometimes you may have more processes than variables.
- Draw a bar graph where the x-axis is the number of processes and the y-axis is the time. There must be 4 bars per x data point on the same graph.
- What are your conclusions?
- Draw another bar graph that shows the speedup (again 4 bars per data data point are expected).
- What are your conclusions?
- Include the two graphs and your conclusions in a single pdf file that has your name and titled *lab1 submission*

What to submit?

Add the source code as well as the pdf file that contains your graphs and conclusions to a zip file named: lastname.firstname.zip

Where lastname is your last name, and firstname is your first name.

How to submit?

email the above zip file to our grader (not to me, not to the mailing list!) with subject line:

lab1 submission

How will we grade this?

- We will test your code with 3 groups of linear equations of different number of variables. The grading will be as follows for each group:
 - Correctness: 5
 - Scalability (we will run your code with increasing number of processes depending on the number of equations): 10
 - This makes a total of 45 points.
- Your graphs and conclusions: 10 points
- Style (i.e. comments, spaces, ...): 5 points.

This makes a total of 60 points: $(3 \cdot 15) + 10 + 5$

Penalties

- If you do not follow the above protocol for submission and file name --> -1
- If your code does not compile, we will look at your source code; and out of the 45 points dedicated to the coding you may not get more than 30 (depending on how close your code is to a correct version).
- You will lose points also if your conclusions are like "As we can see from the graphs, x increasing with y". We need your explanation of what the graphs show, not your description of what we already see!