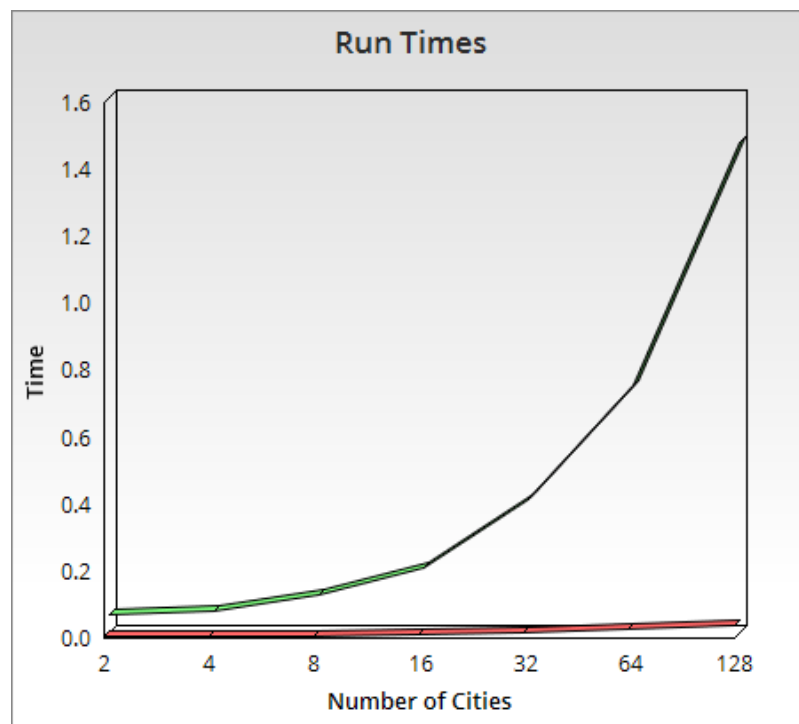


Parallel Computing Lab 2

Andrew Gaffney

Num City	Seq	Parallel
2	.002	.051
4	.004	.064
8	.005	.111
16	.008	.189
32	.011	.396
64	.023	.736
128	.035	1.46



Conclusions

I implemented the traveling salesman problem with a simple nearest neighbor heuristic, that, while quite efficient sequentially, rarely finds the

optimal solution. Additionally, it proved to be quite poor in parallel. Though I wish the green line in the graph above was the sequential performance and the red, the parallel, it is sadly not the case. Since the code contained only one parallelizable for loop, the overhead of thread creation at run-time proved to be quite severe, significantly increasing the run-time over the sequential version.